
Building AI, Using AI



Workshop @ KCDC 2025

~ Shashank Daté ~

GitHub: /shanko

stdate@gmail.com

<https://www.linkedin.com/in/shashankdate/>

Attention is all you need *

* <https://arxiv.org/abs/1706.03762> (Jun 12, 2017)

Agenda:

There will be four sub-sessions, about 1 hour each:

- 1. Intro and Set up:** Intro to various AI/ML terms and prepare your device with necessary tools needed for the workshop.
- 2. AI Coding exercises:** Using AI coding tools to develop scripts consuming various local and cloud based APIs
- 3. Building a RAG application:** A fast paced tutorial on how to build a Retrieval-Augmented Generation model, using AI generated code
- ~~4. Getting started with Fine tuning: learn how to configure the model and the training loop on a resource rich cloud platform~~
- 4. Intro to Agentic AI:** Build a workflow based on a couple of AI Agents

Set Up

Prerequisites: Laptop with at least **16 GB of RAM** and **40 GB of free disk space**

1. Signup for a Github Account: <https://github.com> and clone the repo:
2. If you are on Windows:
 - a. install WSL2: <https://learn.microsoft.com/en-us/windows/wsl/install>
 - b. then follow the remaining instructions on the installed Linux distribution
 - c. OR install Ollama for [Windows](#) (new)
3. Install Python: <https://python.org> version 3.10 or above
4. Create and activate virtual env for python3 that you installed, in the cloned repo

```
$ python3 -m venv venv
```

```
$ source venv/bin/activate
```

5. Install the following Python packages in that environment using pip:

```
$ pip install langchain langchain-chroma langchain-community langchain-core  
langchain-ollama langgraph ollama pandas torch transformers
```

6. Install Ollama: <https://ollama.com>

Optional Set Up

Prerequisites: Please sign the **Liability Waiver** document

- Install Windsurf: <https://windsurf.com>
- Install ngrok: <https://ngrok.com/downloads/mac-os>
- Signup for Lovable: <https://lovable.dev>
- Signup for Google Colab: <https://colab.research.google.com>
- Signup for a Huggingface: <https://huggingface.co>

Start with ChatGPT

Prerequisites: Open ChatGPT: <https://chatgpt.com/>

Prompts:

- What is ChatGPT?
- What is LLM?
- What makes an LLM qualify as “Large”?
- Are there medium, small, tiny LMs?
- What is the LLM on which you are based?
- What is your knowledge cutoff date?
- What is the size of your context window in terms of English words?
- What is the "temperature" setting for GPT-4o on which you are based?

Basic Concepts

What Is:

- Language Model
- Generative AI
- GPT
- Prompt Engineering
- Vibe Coding
- RAG
- Agentic AI

Intro to AI coding tools: Lovable.dev

Competition: famous.ai, bolt.new, ...

Prompt: Create a Progressive Web App that fetches a list of LLMs from a REST API GET request. The display should have a 3-pane layout with the left pane for Navigation, the central wider pane for Chatbot interface, the right pane for Setting the parameters of the LLM. The API end-point URL will be entered via text box in the Settings pane . When the green “Get” button which is below the text-box is clicked, parse the JSON response that is generated by the GET call to extract the names of LLMs. Show the names of the LLMs in a dropdown. When a user selects a name, display a chatbot interface in the central pane, that interacts with the selected LLM. The UI/UX should be minimalistic, clean and have a button to toggle between light and dark mode.

Intro to AI coding tools: Windsurf.com

Competition: Github Copilot, Cursor, ...

Prompt: Review the entire code in the /src folder and identify functions which can be unit tested. For these functions, write unit-tests in a separate test files, one test file per code file. All these test files must be in /test folder. Modify the [README.md](#) file to include instructions on how to run the tests. Do not change the original code in any other way.

Intro to Ollama

What is it: Ollama is a lightweight, extensible, open-source framework for building and running AI Language Models on the local machine.

Provides:

1. Command-Line Interface
2. REST like API
3. Quantization: reduce the computational load by optimizing model performance
4. Customization: via “model file” - a text file that defines how a model should be built, customized, and configured.
5. Lightweight Finetuning: using a method called LoRA (Low-Rank Adaptation)
6. Distribution: via Huggingface

Ollama Continued

- [Accessing](#) via ngrok:

```
ngrok http 11434 --host-header="localhost:11434"
```

- GET /api/version
 - GET /api/ps
 - GET /api/tags
- [Importing](#) a GGUF based model
- [Quantizing](#) a Model

Intro to Retrieval Augmented Generation (RAG)

What is it: a technique that enhances the capabilities of AI Language Models by combining their generative abilities with information retrieval from external sources such as databases.

How it works:

1. **Retrieval:** When a user asks a question, a RAG system first searches a knowledge base (which could be anything from a database of documents to a company's internal knowledge base) to find relevant information.
2. **Augmentation:** This retrieved information is then combined with the original user query to create an augmented prompt.
3. **Generation:** The LLM then uses this augmented prompt to generate a response, drawing on both its pre-trained knowledge and the retrieved information.

Intro to Agentic AI

What is it: An system of “AI agents” that acts autonomously to accomplish specific goals with minimal or no human supervision, using planning, reasoning, and contextual adaptation.

Key characteristics:

- **Autonomy:** Acts independently to achieve predetermined goals without constant human oversight.
- **Proactive reasoning and planning:** Analyzes context, reasons about options, and develops plans rather than following preset rules.
- **Adaptability:** Alters its approach in response to new information or shifting conditions.
- **Multi-agent orchestration:** Often combines multiple AI models or agents that collaborate to complete complex tasks.

Types of AI Agents

- **Simple Reflex Agent:** needs predefined condition-action rules to which it reacts to, e.g. Thermostat
- **Model Based Reflex Agent:** needs state, which it remembers, in addition to the condition-action rules, e.g. Robotic Vacuum Cleaner
- **Goal Based Agent:** no need to have predefined rules, but has goals and it predicts which actions are needed to achieve them, e.g. Self Driving Cars
- **Utility Based Agent:** needs goals prediction plus evaluates how desirable the outcome is based on the utility (scoring/ranking) e.g. Drones
- **Learning Agent:** has all of the above, but improves by learning about the environment from experience e.g. Chess Bots

Caveat: Human-In-The-Loop is still needed!

Resources

Attention is all you need:

<https://arxiv.org/abs/1706.03762>

Windows WSL2:

<https://learn.microsoft.com/en-us/windows/wsl/install>

Python: <https://python.org>

Ollama: <https://ollama.com>

Windsurf: <https://windsurf.com>

Google Colab: <https://colab.research.google.com>

Lovable: <https://lovable.dev>

Huggingface: <https://huggingface.co>

ChatGPT: <https://chatgpt.com>

LangGraph: <https://www.langchain.com/langgraph>

Lambda Labs: <https://lambda.ai>

Runpod: <https://www.runpod.io>

IBM Risk Atlas:

<https://www.ibm.com/docs/en/watsonx/saas?topic=ai-risk-atlas>

Bad LLMs:

<https://blog.sshh.io/p/how-to-backdoor-large-language-models>



**Titanium
Sponsors**



**Platinum
Sponsors**

**Gold
Sponsors**



Other Awesome Organizations



Thanks!

- Event Organizers: esp. Gabby Spurling
- Aalap Patil: <https://www.linkedin.com/in/aalap-patil/>
- Rothbright: <https://rothbright.com>
- GrowthNet: <https://growthnet.biz>
- 360Security: <https://360security.com>
- Apisdor: <https://www.apisdor.com>

Full Workshop: Discount

<https://sysprotech.com/AIWorkshop.html>

Discount Code: Text 913-322-0779

Feedback, please!

